

NlsyLinks

An R package for research with the NLSY (National Longitudinal Survey of Youth)

William H Beasley · Michael D. Hunter ·
David E. Bard · S. Mason Garrison ·
Patrick O’Keefe · Kelly Meredith ·
Joseph Lee Rodgers ·

Received: date / Accepted: date

Abstract The text of your abstract. 150 – 250 words.

Keywords Behavior Genetics · National Longitudinal Survey of Youth · kinship ·

1 Introduction

– Benefits of Accounting for Kinships

William H Beasley
Howard Live Oak, LLC
E-mail: wibeasley@hotmail.com

Michael D. Hunter
Georgia Institute of Technology, Atlanta

David E. Bard
University of Oklahoma Health Sciences Center, OKC

S. Mason Garrison
Wake Forest University, Winston-Salem

Patrick O’Keefe
Vanderbilt University, Nashville

Kelly Meredith
Oklahoma Department of Health

Joseph Lee Rodgers
Vanderbilt University, Nashville

- BG
- D’Onofrio-type research

1.1 Structural and Topical information

The NlsyLinks package offers two types of datasets: topical and structural. *Topical datasets* contain predictor and outcome variables typically used to test a focused hypotheses. For instance, the NLSY79 Gen2 variables Rqqq.qq and Rqqq.qq are critical when studying the relationship between conduct disorder and menarche (*e.g.*, Rodgers et al, 2015), but are not relevant to many hypotheses outside these fields.

In contrast, variables in *structural datasets* are not typically directly stated in the hypotheses, yet are essential to many NLSY-related investigations including:

- familial relationships (*e.g.*, Subjects 301 and 302 are half-brothers; Subjects 301 and 403 are first-cousins),
- subject characteristics (*e.g.*, Subject 301 is a Native American female; Subject 607 died from heart disease in 2005; Subject 802 is part of the military over-sample), and
- subject-survey characteristics (*e.g.*, Subject 301 was 15 years old for the 1981 Survey; Subject 301 did not respond to the 1990 survey; Subject 702 completed the NLSY-C survey in 1996 and the NLSY-YA survey in 1998).

The NlsyLinks includes small topical datasets which allows the vignettes and examples to be reproducible and more realistic. The structural datasets are intended to be the authoritative representations, and are the product of two NIH grants (for a complete history of the familial relationships, see Rodgers et al., 2016).

1.2 Terminology

Since the late 1960s, the NLSY has provided multiple longitudinal cohorts to researcher and as a field, we have not always been consistent with cohort terminology. The package pertains to multiple generations of the ‘Nlsy79’ and a single generation of the ‘Nlsy97’. Because the NlsyLinks package structures information within and between generations of the NLSY simultaneously, it requires slightly unconventional NLSY terminology to reduce ambiguity.

The “Nlsy79” sample refers to both the original 12,686 participants interviewed in 1979, and their 11,500+ children (termed “Nlsy79 Gen1” and “Nlsy79 Gen2”, respectively). “Nlsy79 Gen1” reflects the original NLSY79 study, while Nlsy Gen2 reflects their children, the NLSY79 Children and Young Adults. Although the NLSY does not interview “Nlsy79 Gen0” (the parents of Gen1) or “Nlsy79 Gen3” (the children of Gen2), it does contain direct and indirect information about them.

The terminology for the “Nlsy97” sample is similar yet simpler than the Nlsy79, because the explicit respondents come from a single generation (in a sense the Nlsy97 Gen1). A few variables reflect Gen0 and Gen2. In contrast to the Nlsy79, the Nlsy97 contains more information about the housemates.

They both are nationally-representative samples.

More specifically, the Nlsy79 Gen2 participants are the biological offspring of the Nlsy79 Gen1 mothers. Until roughly the age of 14, the “children” completed the NLSY-C survey, and then became “young adults” and completed the NLSY-YA survey. Though they are the same respondents, different funding mechanisms and different survey items necessitate the distinction. This cohort is sometimes abbreviated as “NLSY79-C”, “NLSY79C”, “NLSY-C” or “NLSYC”. This packages uses “Nlsy79 Gen2” to refer to participants of this generation, regardless of their age at the time of the survey.

Within our own research team, we have mostly stopped using terms like “NLSY79”, “NLSY79-C” and “NLSY79-YA”, because we conceptualize it as one big sample containing two related generations. In many senses, the responses collected from the second generation can be viewed as outcomes of the first generation. Likewise, the parents in the first generation provide many responses that can be viewed as explanatory variables for the 2nd generation. Depending on your research, there can be big advantages of using one cohort to augment the other. There are also survey items that provide information about the 3rd generation and the 0th generation.

Ambiguous Categories: In some cases, the kinship of relatives can be safely narrowed to two categories, but not to one. The most common scenario involves the *ambiguous siblings* (among Nlsy79 Gen2 siblings); due to the sampling design, we know that the two participants share the same mother, so they are either full siblings ($R=.50$) or half siblings ($R=.25$). For the $qqq\%$ that we could not satisfactorily classify, we categorize them as “ambiguous siblings” and assign them $R=.375$. Another common scenario involves the *ambiguous twins* (among Nlsy97 and both generations of Nlsy79). For $qqq\%$ of the twins we comfortably classify their relationship as either MZ or DZ; the remaining $qqq\%$ percent are classified as “ambiguous twins” and assigned $R=0.75$. Empirically, this approach has been successful in the previous 20 years of our previous research, and is discussed further in (Rodgers, qqq).

The `SubjectTag` variable uniquely identify Nlsy79 subjects when a dataset contains both generations. For Gen2 subjects, the `SubjectTag` is identical to their CID (*i.e.*, C00001.00 -the ID assigned in the NLSY79-Children files). However for Gen1 subjects, the `SubjectTag` is their CaseID (*i.e.*, R00001.00), with “00” appended. This manipulation is necessary to identify subjects uniquely in inter-generational datasets. A Gen1 subject with an ID of 43 becomes 4300. The `SubjectTags` of her four children remain 4301, 4302, 4303, and 4304.

The *expected coefficient of relatedness* of a pair of subjects is typically represented by the statistical variable R . Examples are: Monozygotic twins have $R=1$; dizygotic twins have $R=0.5$; full siblings (*i.e.*, those who share both biological parents) have $R=0.5$; half-siblings (*i.e.*, those who share exactly one

biological parent) have $R=0.25$; adopted siblings have $R=0.0$. Other uncommon possibilities are mentioned in the documentation for `Links79Pair`.

In both the Nlsy79 and Nlsy97 samples, the `ExtendedID` variable indicates a subject's extended family. Because there are two generations within the Nlsy79 sample, two subjects will be in the same extended family if either: [1] they are Gen1 housemates, [2] they are Gen2 siblings, [3] they are Gen2 cousins (*i.e.*, they have mothers who are Gen1 sisters in the NLSY79), [4] they are mother and child (in Gen1 and Gen2, respectively), or [5] they are (aunt|uncle) and (niece|nephew) (in Gen1 and Gen2, respectively). Because there is only one generation in the Nlsy97, two subjects will be in the same extended family only if they are housemates.

Note that housemates range from sharing all their genes (*i.e.*, monozygotic twins) to none (*e.g.*, adopted siblings). A primary product of our efforts has been to leverage implicit and explicit information about the subjects in order to facilitate behavior genetics (BG) research, or any other investigation that needs to control for shared genetics.

An **outcome variable** is directly relevant to the applied researcher; these might represent constructs like height, IQ, and income. A **plumbing variable** is necessary to manage BG datasets; examples are `R`, a subject's ID, and the date of a subject's last survey.

An **ACE model** is the basic biometrical model used by BG researchers, where the genetic and environmental effects are assumed to be additive. The three primary variance components are (1) the proportion of variability due to a shared genetic influence (typically represented as a^2 , or sometimes h^2), (2) the proportion of variability due to shared common environmental influence (typically c^2), and (3) the proportion of variability due to unexplained/residual/error influence (typically e^2).

The variables are scaled so that they account for all observed variability in the outcome variable; specifically: $a^2 + c^2 + e^2 = 1$. Using appropriate designs that can logically distinguish these different components (under carefully specified assumptions), the basic biometrical modeling strategy is to estimate the magnitude of a^2 , c^2 , and e^2 within the context of a particular model. For gentle introductions to BG research, we recommend Plomin (1990) and Carey (2003). For more in-depth ACE model-fitting strategies, we recommend Neale & Maes, (2004).

2 Retrieving Data with the NLS Investigator

When a researcher pursues a new idea, we suggest to start by exploring what the NLSY can offer by poking around the (a) vast online documentation and (b) NLS Investigator. The documentation online (<https://www.nlsinfo.org/content/getting-started>) has general information (*e.g.*, how to connect the nationally representative sample was collected), topical information (*e.g.*, what medical and health information has been collected across survey waves and subject ages), and descriptive summaries (*e.g.*, attrition over time

for different race and ethnic groups). This material has helpful suggestions which variables are available and appropriate.

With these hints, identify and download the specific variables from the NLS Investigator. The NLS Investigator is described briefly in the second example. For more detailed instruction see the NlsyLink package's NLS Investigator vignette or the official NLS documentation). Researchers new to the NLSY should expect at least a dozen round trips as they iteratively improve and complete their set of variables. First, select the 'Study', such as 'NLSY79 Child & Young Adult' (which corresponds to 'Nlsy79 Gen2' in our terminology. Second, select your desired variables, out of the tens of thousands available ones.

Before starting the examples, first verify that the NlsyLinks and dplyr packages are installed correctly; for the later examples, the lavaan package is required. If you are using an older version of R, replace the native pipes (`|>`) with pipes from the magrittr package (`%>%`).

```
# Should evaluate to TRUE.
any(.packages(all.available = TRUE) == "NlsyLinks")
any(.packages(all.available = TRUE) == "dplyr")
any(.packages(all.available = TRUE) == "lavaan")

library(NlsyLinks) # Load the package into the current session.
```

3 Illustrative Examples

We will start with a simple example and gradually increase complexity.

3.1 ACE DF Analysis of One Generation (Nlsy79 Gen1)

The first example should provide an initial feel for the inputs, mechanics, and goals of a typical NLSY analysis. It uses a simple statistical model and all available Nlsy79 Gen1 subjects. The `CreatePairLinksDoubleEntered` function will create a data frame where each represents one pair of siblings, respective of order (*i.e.*, there is a row for Subjects 201 and 202, and a second row for Subjects 202 and 201). This function examines the subjects' IDs and determines who is related to whom (and by how much). By default, each row has at least six values/columns: (i) ID for the older member of the kinship pair: `Subject1Tag`, (ii) ID for the younger member: `Subject2Tag`, (iii) ID for their extended family: `ExtendedID`, (iv) their estimated coefficient of genetic relatedness: `R`, (v *and beyond*) outcome values for the older member; (vi *and beyond*) outcome values for the younger member.

A DeFries-Fulker (**DF**) Analysis uses linear regression to estimate the a^2 , c^2 , and e^2 of a univariate biometric system. The interpretations of the DF analysis can be found in Rodgers & Kohler (2005) and Rodgers, Rowe, & Li

(1999). This example uses the newest variation, which estimates two parameters; the corresponding function is called `DeFriesFulkerMethod3`. The steps are:

1. Use the NLS Investigator to select and download a Gen2 dataset.
2. Open R and create a new script (see Appendix: R Scripts and load the NlsyLinks package. If you haven't done so, install the NlsyLinks package). Within the R script, identify the locations of the downloaded data file, and load it into a data frame.
3. Within the R script, load the linking dataset. Then select only Gen2 subjects. The 'Pair' version of the linking dataset is essentially an upper triangle of a symmetric sparse matrix.
4. Load and assign the `ExtraOutcomes79` dataset.
5. Specify the outcome variable name and filter out all subjects who have a negative value in this variable. The NLSY typically uses negative values to indicate different types of missingness (see 'Further Information' below).
6. Create a double-entered file by calling the 'CreatePairLinksDoubleEntered' function. At minimum, pass the (i) outcome dataset, the (ii) linking dataset, and the (iii) name(s) of the outcome variable(s). *(There are occasions when a single-entered file is more appropriate for a DF analysis. See Rodgers & Kohler, 2005, for additional information.)*
7. Use 'DeFriesFulkerMethod3' function (*i.e.*, general linear model) to estimate the coefficients of the DF model.

```
# Step 2: Load the package containing the linking routines.
library(NlsyLinks)
```

```
# Step 3: Load the LINKING dataset and filter for the Gen2 subjects
dsLinking <-
  Links79Pair |>
  dplyr::filter(RelationshipPath == "Gen2Siblings")
summary(dsLinking) # Notice there are 11k records (one for each pair).
```

```
##      ExtendedID      SubjectTag_S1      SubjectTag_S2      R
## Min.      :    2      Min.      :   201      Min.      :   202      Min.      :0.250
## 1st Qu.: 3154      1st Qu.: 315401      1st Qu.: 315403      1st Qu.:0.250
## Median : 6105      Median : 610703      Median : 610705      Median :0.500
## Mean    : 5927      Mean    : 593061      Mean    : 593063      Mean    :0.417
## 3rd Qu.: 8507      3rd Qu.: 851001      3rd Qu.: 851003      3rd Qu.:0.500
## Max.    :12673      Max.    :1267301      Max.    :1267302      Max.    :1.000
##      RelationshipPath
## Gen1Housemates:    0
## Gen2Siblings   :11114
## Gen2Cousins    :    0
## ParentChild    :    0
## AuntNiece      :    0
##
```

Step 4: Load the OUTCOMES dataset, and then examine the summary.

```
dsOutcomes <- ExtraOutcomes79 #'ds' stands for 'Data Set'
summary(dsOutcomes)
```

```
##      SubjectTag      SubjectID      Generation      HeightZGenderAge
## Min.      :   100      Min.      :    1      Min.      :1.000      Min.      : -2.985
## 1st Qu.: 314025      1st Qu.:   5998      1st Qu.:1.000      1st Qu.: -0.724
## Median : 620050      Median :   12000      Median :1.000      Median : -0.045
## Mean    : 618600      Mean    : 289254      Mean    :1.476      Mean    : -0.006
## 3rd Qu.: 914501      3rd Qu.: 577403      3rd Qu.:2.000      3rd Qu.:  0.648
## Max.    :1268600      Max.    :1267501      Max.    :2.000      Max.    :  2.996
##                                     NA's      :4711
## WeightZGenderAge AfqtRescaled2006Gaussified      Afi      Afm
## Min.      : -2.985      Min.      : -2.895      Min.      :  2.00      Min.      :  0.00
## 1st Qu.: -0.677      1st Qu.: -0.692      1st Qu.:15.00      1st Qu.:12.00
## Median : -0.149      Median : -0.024      Median :17.00      Median :13.00
## Mean     :  0.001      Mean     : -0.011      Mean     :16.66      Mean     :12.78
## 3rd Qu.:  0.533      3rd Qu.:  0.660      3rd Qu.:18.00      3rd Qu.:14.00
## Max.     :  4.945      Max.     :  2.994      Max.     :27.00      Max.     :19.00
## NA's      :4719      NA's      :12510      NA's      :12740      NA's      :18165
## MathStandardized
## Min.      : 65.0
## 1st Qu.: 92.5
## Median :100.0
## Mean     :100.1
## 3rd Qu.:108.5
## Max.     :135.0
## NA's      :15085
```

Step 5: If the negative values (which represent NLSY missing or skip patterns) still exist, then:

```
dsOutcomes <-
  dsOutcomes |>
  dplyr::mutate(
    MathStandardized = dplyr::if_else(
      MathStandardized < 0,
      NA_real_,
      MathStandardized
    ),
  )
```

Step 6: Create the double entered dataset.

```
dsDouble <-
  CreatePairLinksDoubleEntered(
    outcomeDataset = dsOutcomes,
    linksPairDataset = dsLinking,
```

```

    outcomeNames      = c('MathStandardized')
  )
  # Notice there are now two records for each unique pair.
  summary(dsDouble)

## SubjectTag_S1      SubjectTag_S2      ExtendedID      R
## Min.      :    201  Min.      :    201  Min.      :    2  Min.      :0.250
## 1st Qu.: 315402  1st Qu.: 315402  1st Qu.: 3154  1st Qu.:0.250
## Median : 610704  Median : 610704  Median : 6105  Median :0.500
## Mean   : 593062  Mean   : 593062  Mean   : 5927  Mean   :0.417
## 3rd Qu.: 851002  3rd Qu.: 851002  3rd Qu.: 8508  3rd Qu.:0.500
## Max.   :1267302  Max.   :1267302  Max.   :12673  Max.   :1.000
##
##      RelationshipPath MathStandardized_S1 MathStandardized_S2
## Gen1Housemates:      0  Min.      : 65.0      Min.      : 65.0
## Gen2Siblings :22228  1st Qu.: 90.0      1st Qu.: 90.0
## Gen2Cousins  :      0  Median : 98.5      Median : 98.5
## ParentChild  :      0  Mean   : 98.4      Mean   : 98.4
## AuntNiece    :      0  3rd Qu.:107.0     3rd Qu.:107.0
##              Max.   :135.0      Max.   :135.0
##              NA's   :3926      NA's   :3926

# Step 7: Estimate the ACE components with a DF Analysis
ace <-
  DeFriesFulkerMethod3(
    dataSet = dsDouble,
    oName_S1 = "MathStandardized_S1",
    oName_S2 = "MathStandardized_S2"
  )
ace

## [1] "Results of ACE estimation: [show]"
##      ASquared      CSquared      ESquared      CaseCount
## 7.734447e-01 1.469204e-01 7.963486e-02 1.668000e+04

```

Further Information: If the different reasons of missingness are important, further work is necessary. For instance, some analyses that use item Y19940000 might need to distinguish a response of “Don’t Know” (which is coded as -2) from “Missing” (which is coded as -7). For this example, it is adequate to combine the non-response values.

3.2 ACE DF Analysis of One Generation (Nlsy79 Gen2)

The second example differs from the previous example in two ways. First, the outcome variables are read from a CSV (comma separated values file) that

was downloaded from the NLS Investigator. Second, the DF analysis is called through the function `AceUnivariate`; this function is a wrapper around some simple ACE methods, and will help us smoothly transition to more techniques later.

The steps are:

1. Use the NLS Investigator to select and download a Gen2 dataset. Select the variables ‘length of gestation of child in weeks’ (C03280.00), ‘weight of child at birth in ounces’ (C03286.00), and ‘length of child at birth’ (C03288.00), and then download the *.zip file to your local computer.
2. Open R and create a new script and load the NlsyLinks package.
3. Within the R script, load the linking dataset. Then select only Gen2 subjects.
4. Read the CSV into R as a `data.frame` using `ReadCsvNlsy79Gen2`.
5. Verify the desired outcome column exists, and rename it something meaningful to your project. In this example, we rename column C0328800 to `BirthWeightInOunces`.
6. Filter out all subjects who have a negative `BirthWeightInOunces` value. See the ‘Further Information’ note in the previous example.
7. Create a double-entered file by calling the `CreatePairLinksDoubleEntered` function. At minimum, pass the (i) outcome dataset, the (ii) linking dataset, and the (iii) name(s) of the outcome variable(s).
8. Call the `AceUnivariate` function to estimate the coefficients.

```
# Step 2: Load the package containing the linking routines.
library(NlsyLinks)

# Step 3: Load the linking dataset and filter for the Gen2 subjects
dsLinking <-
  Links79Pair |>
  dplyr::filter(RelationshipPath == "Gen2Siblings")

# Step 4: Read the outcomes dataset then examine the summary.
# Your path might be `filePathOutcomes <- "C:/Nls/gen2-birth.csv"`
filePathOutcomes <-
  system.file(
    "extdata/gen2-birth.csv",
    package = "NlsyLinks"
  )
dsOutcomes <- ReadCsvNlsy79Gen2(filePathOutcomes)
summary(dsOutcomes)
```

##	SubjectTag	SubjectID	ExtendedID	Generation
##	Min. : 201	Min. : 201	Min. : 2	Min. :2
##	1st Qu.: 310302	1st Qu.: 310302	1st Qu.: 3101	1st Qu.:2
##	Median : 604607	Median : 604607	Median : 6045	Median :2
##	Mean : 601313	Mean : 601313	Mean : 6007	Mean :2

```
## 3rd Qu.: 876203 3rd Qu.: 876203 3rd Qu.: 8757 3rd Qu.:2
## Max. :1267501 Max. :1267501 Max. :12675 Max. :2
## NA's :2
## SubjectTagOfMother C0005300 C0005400 C0005700
## Min. : 200 Min. :1.000 Min. : -3.000 Min. : -3
## 1st Qu.: 310300 1st Qu.:2.000 1st Qu.: 1.000 1st Qu.:1981
## Median : 604600 Median :3.000 Median : 1.000 Median :1985
## Mean : 601311 Mean :2.338 Mean : 1.489 Mean :1986
## 3rd Qu.: 876200 3rd Qu.:3.000 3rd Qu.: 2.000 3rd Qu.:1990
## Max. :1267500 Max. :3.000 Max. : 2.000 Max. :2008
##
## C0328000 C0328600 C0328800
## Min. : -7.00 Min. : -7.0 Min. : -7.00
## 1st Qu.:37.00 1st Qu.: 99.0 1st Qu.:18.00
## Median :39.00 Median :115.0 Median :20.00
## Mean :33.51 Mean :103.9 Mean :16.51
## 3rd Qu.:39.00 3rd Qu.:128.0 3rd Qu.:21.00
## Max. :51.00 Max. :768.0 Max. :48.00
##
```

Step 5: Verify and rename an existing column.

`VerifyColumnExists(dsOutcomes, "C0328600")` *# Should return '10' in this example.*

```
## [1] 10
```

```
dsOutcomes <-
  dsOutcomes |>
  dplyr::rename(
    BirthWeightInOunces = C0328600,
  )
```

*# Step 6: For this item, a negative value indicates the parent
refused, didn't know, invalidly skipped, or was missing.
For present purposes, treat these responses equivalently.
Then clip/Winsorize/truncate the weight.*

```
dsOutcomes <-
  dsOutcomes |>
  dplyr::mutate(
    # Set to NA
    BirthWeightInOunces = dplyr::if_else(
      BirthWeightInOunces < 0L,
      NA_integer_,
      BirthWeightInOunces
    ),
    # Winsorize
    BirthWeightInOunces = pmin(BirthWeightInOunces, 200L),
```

```

)

# Step 7: Create the double entered dataset.
dsDouble <-
  CreatePairLinksDoubleEntered(
    outcomeDataset = dsOutcomes,
    linksPairDataset = dsLinking,
    outcomeNames = c("BirthWeightInOunces")
  )

# Step 8: Estimate the ACE components with a DF Analysis
ace <-
  AceUnivariate(
    method = "DeFriesFulkerMethod3",
    dataSet = dsDouble,
    oName_S1 = "BirthWeightInOunces_S1",
    oName_S2 = "BirthWeightInOunces_S2"
  )
ace

## [1] "Results of ACE estimation: [show]"
##      ASquared      CSquared      ESquared      CaseCount
## 5.103108e-01 1.752415e-01 3.144477e-01 1.744000e+04

```

For another example of incorporating CSVs downloaded from the NLS Investigator, please see the “Race and Gender Variables” entry in the FAQ.

3.3 ACE SEM of One Generation (Nlsy79 Gen2)

The example differs from the first one by the statistical mechanism used to estimate the components. The first example uses multiple regression to estimate the influence of the shared genetic and environmental factors, while this example uses structural equation modeling (SEM).

The `CreatePairLinksSingleEntered` function will create a `data.frame` where each row represents one unique pair of siblings, *irrespective of order*. Other than producing half the number of rows, this function is identical to `CreatePairLinksDoubleEntered`.

The steps are:

(Steps 1-5 proceed identically to the first example.)

6. Create a *single*-entered file by calling the `CreatePairLinksSingleEntered` function. At minimum, pass the (i) outcome dataset, the (ii) linking dataset, and the (iii) name(s) of the outcome variable(s).
7. Declare the names of the outcome variables corresponding to the two members in each pair. Assuming the variable is called ‘ZZZ’ and the preceding

- steps have been followed, the variable 'ZZZ_S1' corresponds to the first members and ZZZ_S2' corresponds to the second members.
8. Create a GroupSummary `data.frame`, which identifies the R groups that should be considered by the model. Inspect the output to see if the groups show unexpected or fishy differences.
 9. Create a `data.frame` with cleaned variables to pass to the SEM function. This `data.frame` contains only the three necessary rows and columns.
 10. Estimate the SEM with the lavaan package. The function returns an `S4` object, which shows the basic ACE information.
 11. Inspect details of the SEM, beyond the ACE components. In this example, we look at the fit stats and the parameter estimates. The lavaan package has additional methods that may be useful for your purposes.

```
# Steps 1-5 are explained in the first example:
library(NlsyLinks)
dsLinking <-
  Links79Pair |>
  dplyr::filter(RelationshipPath == "Gen2Siblings")

dsOutcomes <-
  ExtraOutcomes79 |>
  dplyr::mutate(
    MathStandardized = dplyr::if_else(
      MathStandardized < 0,
      NA_real_,
      MathStandardized
    ),
  )

# Step 6: Create the single entered dataset.
dsSingle <-
  CreatePairLinksSingleEntered(
    outcomeDataset = dsOutcomes,
    linksPairDataset = dsLinking,
    outcomeNames = c('MathStandardized')
  )

# Step 7: Declare the names for the two outcome variables.
oName_S1 <- "MathStandardized_S1" #Stands for Outcome1
oName_S2 <- "MathStandardized_S2" #Stands for Outcome2

# Step 8: Summarize the R groups and
# determine which groups can be estimated.
dsGroupSummary <- RGroupSummary(dsSingle, oName_S1, oName_S2)
dsGroupSummary
```

```
##      R Included PairCount      01Mean      02Mean 01Variance 02Variance
## 1 0.250      TRUE      2689  95.10450  95.97936   126.9489   150.1775
## 2 0.375      TRUE      137  93.63139  93.36861   160.0120   136.6628
## 3 0.500      TRUE     5491  99.89374 100.02868   168.7326   172.7293
## 4 0.750     FALSE       2 108.50000 106.00000   220.5000    18.0000
## 5 1.000      TRUE       21  98.21429  96.02381   289.4393   215.2369
##      0102Covariance Correlation Determinant PosDefinite
## 1      41.96914    0.3039577   17303.459          TRUE
## 2      50.39790    0.3408090   19327.735          TRUE
## 3      90.04116    0.5274225   21037.642          TRUE
## 4      63.00000    1.0000000         0.000         FALSE
## 5     229.10714    0.9179130    9807.933          TRUE
```

```
# Step 9: Create a cleaned dataset
```

```
dsClean <-
  CleanSemAceDataset(
    dsDirty      = dsSingle,
    dsGroupSummary = dsGroupSummary,
    oName_S1     = oName_S1,
    oName_S2     = oName_S2
  )
```

```
# Step 10: Run the model
```

```
ace <- AceLavaanGroup(dsClean)
ace
```

```
## [1] "Results of ACE estimation: [show]"
##      ASquared      CSquared      ESquared      CaseCount
##      0.6219253    0.2097338    0.1683408 8338.0000000
```

```
# Notice the 'CaseCount' is 8.5k instead of 17k.
# This is because (a) one pair with R=.75 was excluded, and
# (b) the SEM uses a single-entered dataset (not double).
```

```
# Step 11: Inspect the output further
```

```
library(lavaan)
```

```
## This is lavaan 0.6-9
## lavaan is FREE software! Please report any bugs.
```

```
GetDetails(ace)
```

```
## lavaan 0.6-9 ended normally after 54 iterations
##
##      Estimator                      ML
##      Optimization method          NLMINB
```

```
##      Number of model parameters                32
##      Number of equality constraints             28
##
##      Number of observations per group:
##      1                                2689
##      2                                137
##      3                                5491
##      4                                21
##
## Model Test User Model:
##
##      Test statistic                447.241
##      Degrees of freedom            16
##      P-value (Chi-square)          0.000
##      Test statistic for each group:
##      1                281.866
##      2                30.277
##      3                127.671
##      4                7.428
```

```
# Examine fit stats like Chi-Squared, RMSEA, CFI, etc.
fitMeasures(GetDetails(ace)) # lavaan defines fitMeasures
```

```
##      npar      fmin      chisq      df
##      4.000      0.027      447.241    16.000
##      pvalue    baseline.chisq    baseline.df    baseline.pvalue
##      0.000      2106.324      4.000      0.000
##      cfi      tli      nnfi      rfi
##      0.795      0.949      0.949      NA
##      nfi      pnfi      ifi      rni
##      NA      3.151      0.794      0.795
##      logl    unrestricted.logl      aic      bic
##      -65103.779    -64880.158      130215.557      130243.671
##      ntotal      bic2      rmsea      rmsea.ci.lower
##      8338.000      130230.960      0.114      0.105
##      rmsea.ci.upper    rmsea.pvalue      rmr      rmr_nomean
##      0.123      0.000      9.992      12.765
##      srmr      srmr_bentler    srmr_bentler_nomean      crmr
##      0.130      0.130      0.089      0.178
##      crmr_nomean      srmr_mplus    srmr_mplus_nomean      cn_05
##      0.026      0.153      0.083      491.246
##      cn_01      gfi      agfi      pgfi
##      597.581      0.999      0.999      0.799
##      mfi
##      0.974
```

Table 1 Count of Nlsy79 relationships, by 'RelationshipPath'. (Recall that 'AuntNiece' also contains uncles and nephews.)

RelationshipPath	Relationship Frequency
Gen1Housemates	5,302
Gen2Siblings	11,114
Gen2Cousins	5,000
ParentChild	11,521
AuntNiece	9,899

```
# Examine low-level details like each group's individual estimates
# and SEs. Uncomment next line to view entire output (~4 pages).
# summary(GetDetails(ace))
```

3.4 ACE SEM of Two Generations

{This is a more moderately difficult analysis with a more common estimation mechanism.}

{Benefits of cross-generational analysis}

The example differs from the previous example in one way –all possible pairs are considered for the analysis. Pairs are only excluded (a) if they belong to one of the small R groups that are difficult to estimate, or (b) if the value for adult height is missing. This includes all 42,836 relationships in the follow five types of NLSY79 relationships.

```
Links79Pair |>
  dplyr::count(RelationshipPath, name = "Relationship Frequency")
```

In our opinion, using the intergenerational links is one of the most exciting new opportunities for NLSY researchers to pursue. We will be happy to facilitate such research through consult or collaboration, or even by generating new data structures that may be of value. The complete kinship linking file facilitates many different kinds of cross-generational research, using both biometrical and other kinds of modeling methods.

```
# Steps 1-5 are explained in the vignette's first example:
library(NlsyLinks)
desiredPaths <-
  c(
    "Gen1Housemates",
    "Gen2Siblings", "Gen2Cousins",
    "ParentChild", "AuntNiece"
  )
dsLinking <-
  Links79Pair |>
```

```
dplyr::filter(RelationshipPath %in% desiredPaths)

# Because all five paths are specified, the line above is equivalent to:
# dsLinking <- Links79Pair

dsOutcomes <- ExtraOutcomes79
# The HeightZGenderAge variable is already groomed

# Step 6: Create the single entered dataset.
dsSingle <-
  CreatePairLinksSingleEntered(
    outcomeDataset = dsOutcomes,
    linksPairDataset = dsLinking,
    outcomeNames = c('HeightZGenderAge')
  )

# Step 7: Declare the names for the two outcome variables.
oName_S1 <- "HeightZGenderAge.S1"
oName_S2 <- "HeightZGenderAge.S2"

# Step 8: Summarize the R groups and determine which groups can be estimated.
dsGroupSummary <- RGroupSummary(dsSingle, oName_S1, oName_S2)
dsGroupSummary
```

##	R	Included	PairCount	O1Mean	O2Mean	O1Variance	O2Variance
## 1	0.0625	TRUE	202	0.22968753	-0.07575395	1.0509023	0.8271487
## 2	0.1250	TRUE	2422	-0.02213186	0.01502572	1.0334694	0.9809845
## 3	0.2500	TRUE	7136	-0.05441460	-0.03885252	1.0189426	1.0294526
## 4	0.3750	TRUE	48	0.13536238	-0.10786196	1.1497643	0.9269109
## 5	0.5000	TRUE	14862	-0.05738494	-0.01738686	0.9602559	0.9851045
## 6	0.7500	FALSE	0	NA	NA	NA	NA
## 7	1.0000	TRUE	27	-0.11092113	-0.12601206	0.6934418	0.9891146

##	O1O2Covariance	Correlation	Determinant	PosDefinite
## 1	0.1112158	0.1192871	0.8568836	TRUE
## 2	0.1487539	0.1477368	0.9916897	TRUE
## 3	0.2721042	0.2656790	0.9749124	TRUE
## 4	0.4074779	0.3947123	0.8996908	TRUE
## 5	0.4146816	0.4263636	0.7739915	TRUE
## 6	NA	NA	NA	FALSE
## 7	0.6863894	0.8287857	0.2147630	TRUE

```
# Step 9: Create a cleaned dataset
dsClean <-
  CleanSemAceDataset(
    dsDirty = dsSingle,
    dsGroupSummary = dsGroupSummary,
```



```

    oName_S1      = oName_S1,
    oName_S2      = oName_S2
  )

```

```

# Step 10: Run the model

```

```

ace <- AceLavaanGroup(dsClean)
ace

```

```

## [1] "Results of ACE estimation: [show]"
##      ASquared      CSquared      ESquared      CaseCount
## 7.366543e-01 6.580003e-02 1.975456e-01 2.469700e+04

```

```

# Step 11: Inspect the output further (see the final step two examples above).

```

Notice the ACE estimates are very similar to the previous example, but the number of pairs has increased by 6x –from 4,185 to 24,700. The number of *subjects* doubles when Gen2 is added, and the number of *relationship pairs* increases combinatorially. When an extended family’s entire pedigree is considered by the model, many more types of links are possible than if just nuclear families are considered. This increased statistical power is even more important when the population’s a^2 is small or moderate.

Note that the analysis has 24,697 relationships instead of the entire 42,836. This is primarily because not all subjects have a value for ‘adult height’ (and that’s mostly because a lot of Gen2 subjects are too young). There are 42,088 pairs with a nonmissing value in `RFull1`, meaning that 98.3 are classified. We feel comfortable claiming that if a researcher has a phenotype for both members of a pair, there’s a 99+% chance we have an `RFull1` for it. For a description of the `R` and `RFull1` variables, please see the `Links79Pair` entry in the package reference manual.

3.5 More Advanced ACE Analyses

- When the analysis grows beyond a single outcome at one time point, researchers are better off using the modeling software itself (*e.g.*, `OpenMx`, `lavaan`, `Mplus`) than the wrappers provided by `NlsyLinks`, or any other package.

4 Further Topics

4.1 Data Manipulation and Non-Biometric Analyses

- Even if the investigation doesn’t involve family structure, `NlsyLinks` functions and dataset can make the research can be more efficient.

4.2 SAS Analogues

- The `NlsyLinks` datasets can be used in any statistical package, and we demonstrate that here with SAS.
- Downloadable from `qqq.-qqqq`

4.3 Additional Resources

4.4 References